



UNSW
SYDNEY

SCHOOL OF MATHEMATICS AND STATISTICS

MATH5836 DATA MINING

Course Instructor: Dr. Rohitash Chandra, Office: RC-4110

Email: rohitash.chandra@unsw.edu.au

2020 Units of Credit: 6, Complementary Courses: MATH5855; MATH5856; COMP9417

Prerequisites: A prerequisite for this course is MATH2801/2901. A recommended prerequisite is MATH2831/2931 (Higher Linear Models). If you have not done the recommended prerequisite courses, make sure that you know sufficient statistical theory, can integrate and differentiate, and have competency in at least a programming language such as R or Python.

Delivery Mode: The course will have 4 contact hours of Lecture that will feature tutorials blended in with Lectures. Monday 5-7pm and Tuesday 5-7pm, Trimester 3, 2020. The lectures will be given online and a recording will also be provided. The course will use Moodle and Edstem.

DESCRIPTION

Increasingly, organisations need to analyse enormous data sets to determine useful structure in them. In response to this, a range of statistical methods and tools have been developed in recent times to allow accurate and quick analysis of these datasets. This course covers the key techniques in data mining and machine learning together with theoretical background and applications to the financial environment. Case studies will include industry-based data mining projects. We will cover methods such as linear and logistic regression, neural networks, Bayesian neural networks, clustering and dimensionality reduction, trees and forests, ensemble learning, and emerging data mining methods and applications. Emerging data mining and machine learning tools will be used to illustrate the methods in programming environments such as Python and R.

RATIONALE

New ideas and skills are introduced and demonstrated in lectures and through recommended reading of supplementary material such as research papers, then students develop these skills by applying them to specific tasks in assessments. We believe that effective learning is best supported by a climate of inquiry, in which students are actively engaged in the learning process. Hence this course is structured with a strong emphasis on problem-solving tasks. Students are expected to devote the majority of their class and study time to the solving of such tasks. New ideas and skills are first introduced and demonstrated in lectures, and then students develop these skills by applying them to specific tasks in assessments. This course has a major focus on research, inquiry and analytical thinking as well as information literacy. We will also explore capacity and motivation for intellectual development through the solution of both simple and complex mathematical models of problems arising in finance, economics and engineering, and the interpretation and communication of the results.

AIM

This course is expected to give students an understanding of the fundamentals of machine learning and basics of data mining, which is essential for anyone contemplating a career as a professional statistician or data analyst in industries reliant upon such expertise. The student should develop a working knowledge of the statistical and theoretical underpinnings of the topics covered. Given this fundamental statistical understanding of these methodologies this will allow the student to

utilise these techniques with confidence on real world data sets and scenarios. As such the student is expected to develop applied working knowledge of the methodologies covered, largely through practical applications. In addition, students will undertake additional reading of a collection of associated research papers in each topic, to further add context to the methodologies presented during the course. This will enhance the students ability to utilise these techniques to solve real-world problems. It is stressed that this course is aimed at fundamental statistical properties of these methods, it is not a course on application of computer software.

ASSESSMENT

Note that the course assessment can be done either in R or Python. Support for both languages will be given.

1. Quiz (5%)
2. Assignment 1 (20%)
3. Assignment 2 (25%)
4. Final Exam (50%)

SCHEDULE

1. Week 0 Python and R Tutorials (no Lectures)
2. Week 1 Data Processing and Introduction to Data mining
3. Week 2 Logistic Regression and Evaluation
4. Week 3 Intro to Neural Networks
5. Week 4 Advances in Neural Networks
6. Week 5 Bayesian Neural Networks
7. Week 6 Break (no Lectures)
8. Week 7 Trees and Forests
9. Week 8 Ensemble Learning
10. Week 9 Unsupervised Learning and Dimensionality Reduction
11. Week 10 Emerging Topics in Data Mining

LEARNING OUTCOMES

1. Demonstrate an understanding of the fundamentals of machine learning and basics of data mining.
2. Demonstrate a working knowledge of the statistical and theoretical underpinnings of the methods.
3. Demonstrate an applied working knowledge of the methodologies covered with practical assignments.
4. Develop models for solving data mining problems that include clustering, regression, and classification.
5. Build models and apply them to real-world data sets and use evaluation metrics to compare their performance.

RECOMMENDED BOOK

1. Géron. A, 2019, Hands-on Machine Learning with Scikit-Learn, Keras, and TensorFlow : concepts, tools, and techniques to build intelligent systems, O'Reilly, second edition: <https://www.bookshop.unsw.edu.au/details.cgi?ITEMNO=9781492032649> (a copy would be handy but not required)

EXTRA READING MATERIALS (optional)

2. Mitchell. Tom, 1997, Machine Learning, McGraw-Hill. (additional textbook for reference): <https://www.amazon.com.au/Machine-Learning-Thomas-Mitchell/dp/0070428077>
3. Kroese, Botev, Tamire & Vaisman (2020), Data Science and Machine Learning, Chapman & Hall: <https://www.bookshop.unsw.edu.au/details.cgi?ITEMNO=9781138492530>

SPECIAL CONSIDERATION

You can apply for special consideration if illness or other circumstances beyond your control interfere with your assessment performance, to get an extra opportunity to demonstrate your level of performance.

You must make your application online, through the [Special Consideration portal on myUNSW](#). Do not apply to your course teaching staff - they will be notified automatically.

You can read more about special consideration at: <https://student.unsw.edu.au/special-consideration>.

ACADEMIC INTEGRITY

For further information about academic integrity and plagiarism at UNSW go to: <https://student.unsw.edu.au/plagiarism>

For information about acknowledging your sources and referencing go to: <https://student.unsw.edu.au/referencing>. If you are not sure what referencing style to use in this course, you should ask your Lecturer.